

# Optimization algorithms

**10/10 points (100.00%)**

Quiz, 10 questions

**✓ Congratulations! You passed!**[Next Item](#)1 / 1  
points

1.

Which notation would you use to denote the 3rd layer's activations when the input is the 7th example from the 8th minibatch?

 $a^{[8]\{3\}}(7)$  $a^{[8]\{7\}}(3)$  $a^{[3]\{8\}}(7)$ **Correct** $a^{[3]\{7\}}(8)$ 1 / 1  
points

2.

Which of these statements about mini-batch gradient descent do you agree with?



You should implement mini-batch gradient descent without an explicit for-loop over different mini-batches, so that the algorithm processes all mini-batches at the same time (vectorization).



Training one epoch (one pass through the training set) using mini-batch gradient descent is faster than training one epoch using batch gradient descent.



One iteration of mini-batch gradient descent (computing on a single mini-batch) is faster than one iteration of batch gradient descent.

**Correct**



points

## Optimization algorithms

**10/10 points (100.00%)**

Quiz, 10 questions

Why is the best mini-batch size usually not 1 and not  $m$ , but instead something in-between?

☐

If the mini-batch size is  $m$ , you end up with batch gradient descent, which has to process the whole training set before making progress.

**Correct**☐

If the mini-batch size is 1, you end up having to process the entire training set before making any progress.

**Un-selected is correct**☐

If the mini-batch size is  $m$ , you end up with stochastic gradient descent, which is usually slower than mini-batch gradient descent.

**Un-selected is correct**☐

If the mini-batch size is 1, you lose the benefits of vectorization across examples in the mini-batch.

**Correct**1 / 1  
points


4.

Suppose your learning algorithm's cost  $J$ , plotted as a function of the number of iterations, looks like this:

## Optimization algorithms

10/10 points (100.00%)

Quiz, 10 questions



Which of the following do you agree with?

- ☒ If you're using mini-batch gradient descent, this looks acceptable. But if you're using batch gradient descent, something is wrong.



Correct

- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, this looks acceptable.
- ☐ If you're using mini-batch gradient descent, something is wrong. But if you're using batch gradient descent, this looks acceptable.
- ☐ Whether you're using batch gradient descent or mini-batch gradient descent, something is wrong.



1 / 1  
points

5.

Suppose the temperature in Casablanca over the first three days of January are the same:

## Optimization algorithms

10/10 points (100.00%)

Quiz, 10 questions Jan 1st:  $\theta_1 = 10^\circ C$

Jan 2nd:  $\theta_2 10^\circ C$

(We used Fahrenheit in lecture, so will use Celsius here in honor of the metric world.)

Say you use an exponentially weighted average with  $\beta = 0.5$  to track the temperature:  $v_0 = 0$ ,  $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$ . If  $v_2$  is the value computed after day 2 without bias correction, and  $v_2^{corrected}$  is the value you compute with bias correction. What are these values? (You might be able to do this without a calculator, but you don't actually need one. Remember what is bias correction doing.)

☐  $v_2 = 7.5, v_2^{corrected} = 7.5$

☒  $v_2 = 7.5, v_2^{corrected} = 10$



Correct

☐  $v_2 = 10, v_2^{corrected} = 10$

☐  $v_2 = 10, v_2^{corrected} = 7.5$



1 / 1  
points

6.

Which of these is NOT a good learning rate decay scheme? Here,  $t$  is the epoch number.

☐  $\alpha = 0.95^t \alpha_0$

☐  $\alpha = \frac{1}{\sqrt{t}} \alpha_0$

☐  $\alpha = \frac{1}{1+2*t} \alpha_0$

☒  $\alpha = e^t \alpha_0$



Correct



1 / 1  
points

7.

## Optimization algorithms


Quiz, 10 questions

You use an exponentially weighted average on the London temperature dataset. You

use the following to track the temperature:  $v_t = \beta v_{t-1} + (1 - \beta)\theta_t$ . The red line

below was computed using  $\beta = 0.9$ . What would happen to your red curve as you

vary  $\beta$ ? (Check the two that apply)



☐ Decreasing  $\beta$  will shift the red line slightly to the right.



Un-selected is correct

☐ Increasing  $\beta$  will shift the red line slightly to the right.



Correct

True, remember that the red line corresponds to  $\beta = 0.9$ . In lecture we had a green line ( $\beta = 0.98$ ) that is slightly shifted to the right.

☐ Decreasing  $\beta$  will create more oscillation within the red line.



Correct

True, remember that the red line corresponds to  $\beta = 0.9$ . In lecture we had a yellow line ( $\beta = 0.98$ ) that had a lot of oscillations.

☐ Increasing  $\beta$  will create more oscillations within the red line.



Un-selected is correct




points

## Optimization algorithms

10/10 points (100.00%)

Quiz, 10 questions

Consider this figure:



These plots were generated with gradient descent; with gradient descent with momentum ( $\beta = 0.5$ ) and gradient descent with momentum ( $\beta = 0.9$ ). Which curve corresponds to which algorithm?

- ☐ (1) is gradient descent with momentum (small  $\beta$ ). (2) is gradient descent. (3) is gradient descent with momentum (large  $\beta$ )
- ☒ (1) is gradient descent. (2) is gradient descent with momentum (small  $\beta$ ). (3) is gradient descent with momentum (large  $\beta$ )



Correct

- ☐ (1) is gradient descent. (2) is gradient descent with momentum (large  $\beta$ ). (3) is gradient descent with momentum (small  $\beta$ )
- ☐ (1) is gradient descent with momentum (small  $\beta$ ), (2) is gradient descent with momentum (small  $\beta$ ), (3) is gradient descent

1 / 1  
points

9.

Suppose batch gradient descent in a deep network is taking excessively long to find a value of the parameters that achieves a small value for the cost function  $\mathcal{J}(W^{[1]}, b^{[1]}, \dots, W^{[L]}, b^{[L]})$ . Which of the following techniques could help find parameter values that attain a small value for  $\mathcal{J}$ ? (Check all that apply)

Try tuning the learning rate  $\alpha$ 

Correct



Try mini-batch gradient descent



**Correct**

## Optimization algorithms

**10/10 points (100.00%)**

Quiz, 10 questions



Try initializing all the weights to zero

**Un-selected is correct**

Try better random initialization for the weights

**Correct**

Try using Adam

**Correct**1 / 1  
points

10.

Which of the following statements about Adam is False?

We usually use “default” values for the hyperparameters  $\beta_1, \beta_2$  and  $\varepsilon$  in Adam ( $\beta_1 = 0.9, \beta_2 = 0.999, \varepsilon = 10^{-8}$ )

Adam combines the advantages of RMSProp and momentum



Adam should be used with batch gradient computations, not with mini-batches.

**Correct**The learning rate hyperparameter  $\alpha$  in Adam usually needs to be tuned.